

Data profiling with Apache Calcite

DataWorks Summit 2017
SAN JOSE, USA

2017/06/14
Julian Hyde



@julianhyde

SQL
Query planning
Query federation
OLAP
Streaming
Hadoop



ASF member
Original author of Apache Calcite
PMC Apache Arrow, Drill, Eagle, Kylin



Overview

Apache Calcite

Motivating problem: Automatically designing summary tables

What is data profiling?

Naive profiling algorithm

Improving the algorithm using sketches, parallelism, information theory

Applying data profiling to other problems

Apache Calcite



Apache top-level project since October, 2015

Query planning framework

- Relational algebra, rewrite rules
- Cost model & statistics
- Federation via adapters
- Extensible


Packaging

- Library
- Optional SQL parser, JDBC server
- Community-authored rules, adapters




Planning queries

```
select p.productName, count(*) as c
from splunk.splunk as s
      join mysql.products as p
      on s.productId = p.productId
where s.action = 'purchase'
group by p.productName
order by c desc
```



Optimized query

```
select p.productName, count(*) as c
from splunk.splunk as s
      join mysql.products as p
      on s.productId = p.productId
where s.action = 'purchase'
group by p.productName
order by c desc
```



Optimizing queries

Problem

10 TB database, disk with 1 GB/s throughput, and a query that reads 1 TB data.

Solutions

1. **Sequential scan** Query takes 1,000s.
2. **Parallelize** Spread the data over 100 disks in 25 machines. Query takes 10s.
3. **Cache** Keep the data in memory. 2nd query: 10ms. 3rd query: 10s.
4. **Materialize** Summarize the data on disk. All queries: 100ms.
5. **Materialize + cache + adapt** As above, building summaries on demand.

Optimizing data

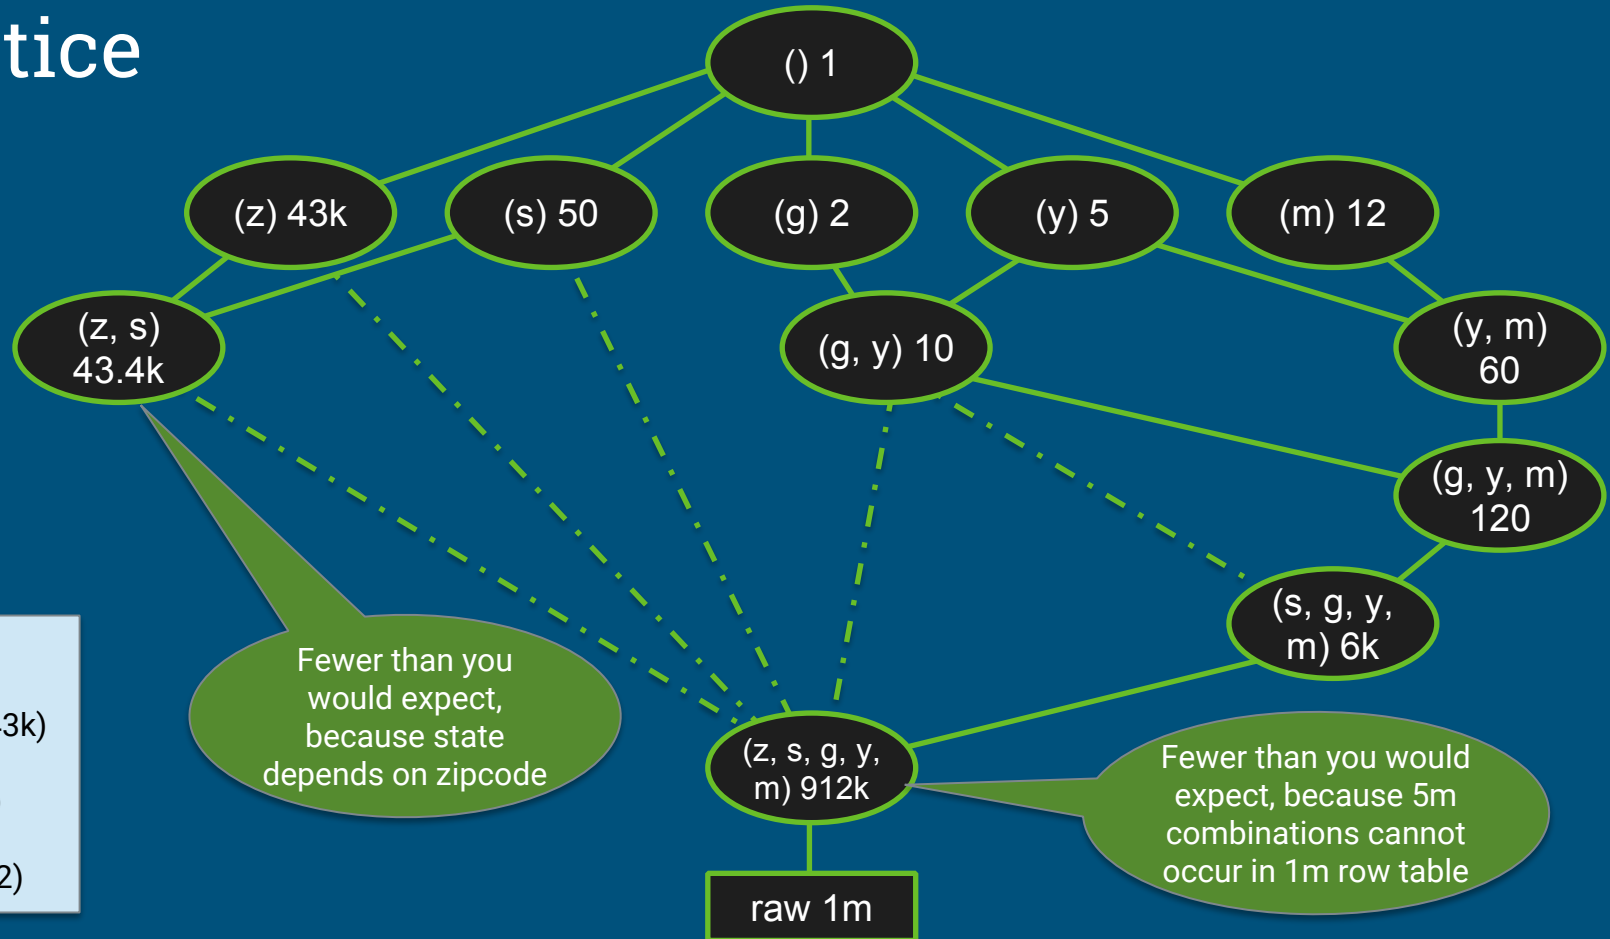
```
create materialized view EmpSummary as
select deptno, COUNT(*) as c, SUM(sal) as s
from Emp
group by deptno
```

A materialized view (“materialization”) is a table that contains the result of a query. The DBMS maintains it, and uses it to answer queries on other tables.

Challenges:

- **Design** Which materializations to create?
- **Populate** Load them with data
- **Maintain** Incrementally populate when data changes
- **Rewrite** Transparently rewrite queries to use materializations
- **Adapt** Design and populate new materializations, drop unused ones
- **Express** Need a rich algebra, to model how data is derived

Lattice



Algorithm: Design summary tables

Given a database with 30 columns, 10M rows. Find X summary tables with under Y rows that improve query response time the most.

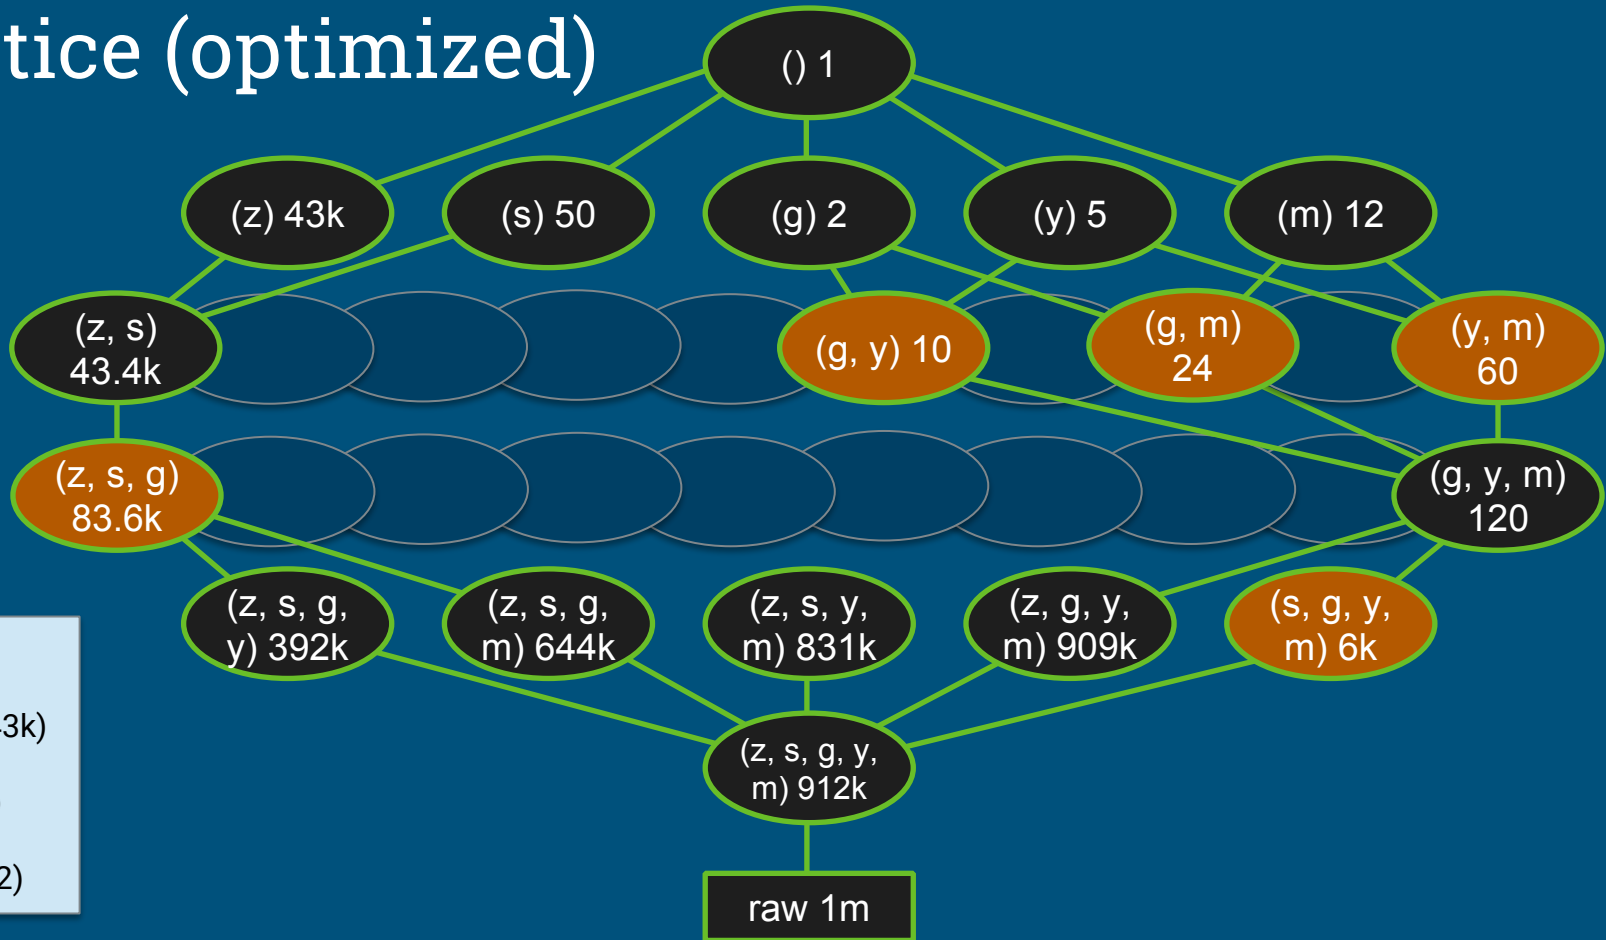
AdaptiveMonteCarlo algorithm [1]:

- Based on research [2]
- Greedy algorithm that takes a combination of summary tables and tries to find the table that yields the greatest cost/benefit improvement
- Models “benefit” of the table as query time saved over simulated query load
- The “cost” of a table is its size

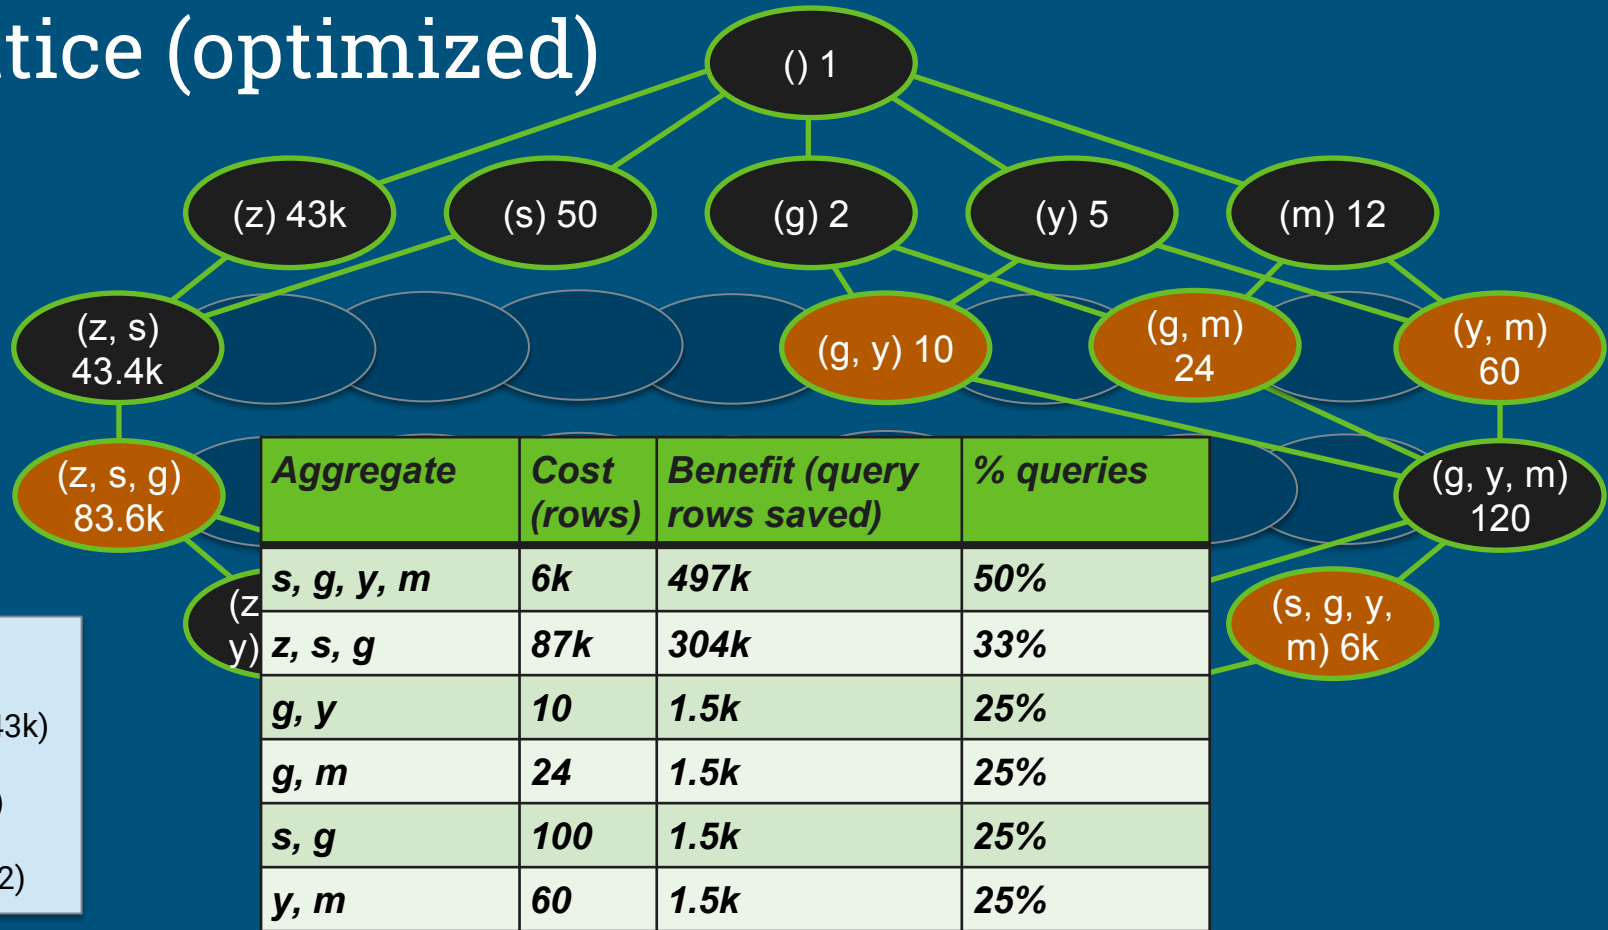
[1] org.pentaho.aggdes.algorithm.impl.AdaptiveMonteCarloAlgorithm

[2] Harinarayan, Rajaraman, Ullman (1996). “Implementing data cubes efficiently”

Lattice (optimized)



Lattice (optimized)



Data profiling

Algorithm needs `count(distinct a, b, ...)` for each combination of attributes:

- Previous example had $2^5 = 32$ possible tables
- Schema with 30 attributes has 2^{30} (about 10^9) possible tables
- Algorithm considers a significant fraction of these
- Approximations are OK

Attempts to solve the profiling problem:


1. Compute each combination: scan, sort, unique, count; repeat 2^{30} times!
2. Sketches (HyperLogLog)
3. Sketches + parallelism + information theory (CALCITE-1616)

Sketches

HyperLogLog is an algorithm that computes approximate distinct count. It can estimate cardinalities of 10^9 with a typical error rate of 2%, using 1.5 kB of memory. [3][4]

With 16 MB memory per machine we can compute 10,000 combinations of attributes each pass.

So, we're down from 10^9 to 10^5 passes.



[3] Flajolet, Fusy, Gandouet, Meunier (2007). "Hyperloglog: The analysis of a near-optimal cardinality estimation algorithm"
[4] <https://github.com/mrjgreen/HyperLogLog>

Combining probability & information theory

Given	Expected cardinality	Actual cardinality	Surprise
(gender): 2 (state): 50	(gender, state): 100.0	100	0.000
(month): 12 (zipcode): 43,000	(month, zipcode): 441,699.3	442,700	0.001
(state): 50 (zipcode): 43,000	(state, zipcode): 799,666.7	43,400	0.897
(state, zipcode): 43,400 (gender, state): 100 (gender, zipcode): 85,995	(gender, state, zipcode): 86,799 = min(86,799, 892,234, 892,228)	83,567	0.019

- Surprise = $\text{abs}(\text{actual} - \text{expected}) / (\text{actual} + \text{expected})$
- $E(\text{card}(x, y)) = n \cdot (1 - ((n - 1) / n)^p)$ $n = \text{card}(x) * \text{card}(y)$, $p = \text{row count}$

Algorithm

Three ways “surprise” can help:


- If a cardinality is not surprising, we don't need to store it -- we can derive it
- If a combination's cardinality is not surprising, it is unlikely to have surprising children
- If we're not seeing surprising results, it's time to stop

```
surprise_threshold := 1
queue := {singleton combinations} // (a), (b), ...
while queue is not empty {
  batch := remove first 10,000 entries in queue
  compute cardinality of each combination in batch
  for each actual (computed) cardinality a {
    e := expected cardinality of combination
    s := surprise(a, e)
    if s > surprise_threshold {
      store combination and its cardinality
      add child combinations to queue // (x, a), (x, b), ...
    }
    increase surprise_threshold
  }
}
```


Algorithm progress and “surprise” threshold

Singleton combinations have surprise = 1

Surprise threshold rises after we have completed the first batch



Surprise threshold rises as algorithm progresses

Rejected as not sufficiently surprising

Progress of algorithm



Hierarchies considered harmful

Hierarchies are a feature of most OLAP systems


Does it makes sense to store (year, quarter, month, date) and roll up to (year, quarter)?

No -- algorithm can deduce hierarchies; less configuration means fewer mistakes


Summary optimizer naturally includes attributes that don't increase summary cardinality by much

Feel free to specify a “drill path” in slice & dice UI

True hierarchy



Almost a hierarchy



Other applications of data profiling

Query optimization:

- Planners are poor at estimating selectivity of conditions after N-way join (especially on real data)
- New join-order benchmark: “Movies made by French directors tend to have French actors”
- Predict number of reducers in MapReduce & Spark

“Grokking” a data set

Identifying problems in normalization, partitioning, quality

Applications in machine learning?

Further improvements

- Build sketches in parallel
- Run algorithm in a distributed framework (Spark or MapReduce)
- Compute histograms
 - For example, Median age for male/female customers
- Seek out functional dependencies
 - Once you know FDs, a lot of cardinalities are no longer “surprising”
 - FDs occur in denormalized tables, e.g. star schemas
- Smarter criteria for stopping algorithm
- Skew/heavy hitters. Are some values much more frequent than others?
- Conditional cardinalities and functional dependencies
 - Does one partition of the data behave differently from others? (e.g. year=2005, state=LA)

Thank you!

<https://issues.apache.org/jira/browse/CALCITE-1788>

<https://calcite.apache.org>

@ApacheCalcite

@julianhyde

